

RAHUL BHATT

Jabalpur, Madhya Pradesh, India | +91 78983 72675 | rahulbhatt.tech@gmail.com

<https://github.com/rahulbhattsd> <https://www.linkedin.com/in/rahulbhatt-developer/>

Professional Summary

AI Engineer focused on **Generative AI, Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), AI agents, semantic search, and LLM-powered application development**. Hands-on experience building retrieval-grounded assistants, AI interview systems, LLM gateways, semantic caching systems, token usage tracking, and real-time AI applications. Experienced with **LangChain, LangGraph, OpenAI API, Anthropic API, Groq API, Python, RAG pipelines, prompt engineering, NLP, and LLM application architecture**.

AI / LLM Skills

Generative AI: Large Language Models (LLMs), Generative AI, AI Agents, Conversational AI, Natural Language Processing (NLP), Prompt Engineering

LLM Frameworks: LangChain, LangGraph

LLM APIs & Models: OpenAI API, Anthropic API, Groq API, Llama 3.3 70B, Multi-Provider LLM Integration

Retrieval-Augmented Generation: RAG, Semantic Search, Retrieval Pipelines, Context Retrieval, Knowledge Grounding, Retrieval-Based Question Answering

AI Agent Engineering: AI Agents, Agent Workflows, Tool Calling, Function Calling, Conversation Memory, Context Management, Workflow Orchestration

LLM Optimization: Semantic Caching, Redis Caching, Token Usage Tracking, LLM Cost Optimization, Rate Limiting, Response Streaming

AI Application Engineering: LLM API Integration, AI Application Architecture, AI Backend Integration, Real-Time AI Applications, Session-Based AI Applications

Supporting Technologies: Python, Node.js, TypeScript, REST APIs, PostgreSQL, MongoDB, Redis, Docker, AWS, Git

Education

Gyan Ganga Institute of Technology and Sciences

Bachelor of Technology (B.Tech), Computer Science and Engineering – CGPA: 7.41

2022 – 2026

Jabalpur, Madhya Pradesh, India

Class XII: 73% – Abbot Baeten School, Jabalpur

Class X: 91% – Abbot Baeten School, Jabalpur

Certifications

- Generative AI with Large Language Models – DeepLearning.AI
- AWS Certified Cloud Practitioner – Amazon Web Services
- Cisco DevNet Associate – Cisco

AI Engineering Projects

TokenGate – LLM Gateway & AI Cost Optimization Platform LLMs, LangChain, OpenAI API, Anthropic API, Node.js, Redis

- Built a centralized **LLM API gateway** for integrating and routing requests across multiple AI model providers through a unified interface.
- Integrated multi-provider LLM workflows supporting centralized model access, request routing, usage tracking, and provider abstraction.
- Implemented **semantic caching with Redis** to identify reusable AI responses and reduce redundant LLM API requests.
- Built **LLM token usage tracking and cost analytics** to monitor model consumption and AI API expenditure.
- Developed request logging and API key management for controlling and monitoring LLM-powered applications.
- Designed the platform as an AI infrastructure layer for applications requiring centralized LLM provider management, caching, usage tracking, and cost optimization.
- **Code:** <https://github.com/rahulbhattsd/TokenGate>

VibeBot – Retrieval-Augmented Generation AI Assistant RAG, Groq API, Llama 3.3 70B, MongoDB, Node.js

- Built a **RAG-powered AI customer-support assistant** using the Groq API and Llama 3.3 70B to generate responses grounded in application data.
- Implemented retrieval of live **product and order information from MongoDB** before LLM generation to provide context-aware responses.
- Designed an AI workflow combining retrieval, contextual grounding, LLM generation, and real-time response delivery.
- Implemented **Server-Sent Events (SSE)** for streaming LLM responses to the client in real time.

- Added rate limiting to control AI API usage and protect the application from excessive requests.
- **Demo:** <https://vibecart-eo6e.onrender.com/> **Code:** <https://github.com/rahulbhatts/VibeCart>

Mock.ai – LLM-Powered Interview Platform

LLMs, LangChain, LangGraph, Generative AI, Node.js

- Built an AI-powered interview platform using **LLM-generated questions** to create role-specific mock interview experiences.
- Designed AI interview workflows that dynamically generate interview content based on the selected role and user interaction.
- Implemented LLM-powered application workflows for interview generation, session management, authentication, and persistent user data.
- Integrated AI functionality into a complete user-facing application with real-time interaction and persistent session data.
- **Demo:** <https://mock-ai-jw8k.onrender.com/> **Code:** <https://github.com/rahulbhatts/Mock.ai>

Healthify – AI Conversational Assistant

Generative AI, LLMs, OpenAI API, Node.js, MongoDB

- Developed an AI conversational application integrating **LLM-powered interactions** with persistent user conversation history.
- Implemented backend services for AI interactions, session management, conversation persistence, and user-specific AI workflows.
- Designed persistent conversation storage in MongoDB to maintain continuity across AI-assisted user sessions.
- Built the complete AI application flow from user interaction and backend orchestration to LLM integration and persistent conversation data.
- **Demo:** <https://healthify-31ok.onrender.com/landing> **Code:** <https://github.com/rahulbhatts/healthify>

Experience

IBM

Jun 2022 – Jul 2022

Front-End Developer Intern – Virtual Internship

Remote

- Developed responsive web interfaces and reusable frontend components during a software development internship.
- Worked with Git-based development workflows, Agile/Scrum practices, testing, and collaborative software development.

Additional Experience

Open Source Contributor

Ongoing

Independent Contributor

GitHub / Remote

- Contributed code, bug fixes, documentation, and engineering improvements to publicly hosted GitHub projects through pull requests and collaborative development workflows.